

# Project 3 Report: User-level Memory Management

CS 416/518: Operating Systems Design

Team Members:

- Kelvin Ihezue (NetID: ki120)
  - Bryan Shangguan (NetID: bys8)
- 

## 1. Detailed Implementation Logic (Virtual Memory Functions)

Refers to the 32-bit implementation in `my_vm.c` and `my_vm.h`.

### Initialization (`set_physical_mem`)

- Allocates a 1GB buffer (`g_physical_mem`) using `malloc` to simulate physical RAM.
- Initializes two bitmaps: `g_phys_bitmap` to track physical frame allocation and `g_virt_bitmap` to track virtual page allocation.
- Allocates the top-level Page Directory.
- Uses `pthread_mutex` locks to ensure thread safety during initialization.

### Address Translation (`translate`)

- **TLB Lookup:** First checks the TLB using `TLB_check`. If valid, returns the cached physical address immediately.
- **Page Walk:** If TLB misses, extracts the Page Directory Index (PDX) and Page Table Index (PTX) using bit masking/shifting macros defined in `my_vm.h`.
- **Two-Level Traversal:**
  - Checks if the PDE is present.
  - Calculates the address of the inner Page Table.
  - Checks if the PTE is present.
- **TLB Update:** If the translation is found in the page table, it adds the entry to the TLB.

### Page Mapping (`map_page`)

- Takes a virtual address and physical address as input.
- Walks the Page Directory. If a Page Table is missing for the required directory entry, it finds a free physical frame, allocates a new Page Table, and updates the PDE.
- Updates the specific Page Table Entry (PTE) with the Physical Frame Number (PFN) and sets status bits (`PTE_PRESENT`, `PTE_WRITABLE`).
- Uses `phys_bitmap_lock` to ensure two threads do not grab the same free frame simultaneously.

## Memory Allocation (`n_malloc`)

- **Size Calculation:** Rounds up the requested bytes to the nearest number of full pages.
- **Virtual Allocation:** Uses `get_next_avail` to scan the virtual bitmap (`g_virt_bitmap`) for a contiguous sequence of free virtual pages.
- **Physical Allocation:** Iterates through the required number of pages. For each, it finds a free physical frame using the physical bitmap.
- **Mapping:** Calls `map_page` to link the virtual page to the physical frame.
- **Error Handling:** If allocation fails at any step (e.g., out of physical memory), it performs a rollback, freeing any bits or pages set during the attempt.

## Memory Deallocation (`n_free`)

- Calculates the number of pages to free based on the size argument.
- **Bitmap Cleanup:** Clears the corresponding bits in `g_virt_bitmap` to release the virtual address space.
- **Page Table Cleanup:** Walks the page table for the specified range. For every valid PTE found:
  - Extracts the PFN.
  - Clears the corresponding bit in `g_phys_bitmap`.
  - Invalidates the PTE (sets it to 0).

## Data Movement (`put_data / get_data`)

- **Boundary Handling:** Handles read/write operations that span multiple pages. It calculates how much data fits in the current page (`PGSIZE - offset`) and copies data in chunks.
- **Translation:** Calls `translate` for every page accessed to ensure the physical address is valid before copying.
- **Memcpy:** Uses `memcpy` to transfer data between the user buffer and the simulated physical memory (`g_physical_mem`).

## TLB Implementation (`add_TLB, check_TLB`)

- **Structure:** Implements a direct-mapped cache.
- **Indexing:** Uses modulo (`vpn % TLB_ENTRIES`) to determine the index.
- **Eviction:** Because it is direct-mapped, collisions simply overwrite the existing entry at that index (compulsory eviction).
- **Miss Rate:** Tracks total lookups and misses to calculate the rate in `print_TLB_missrate`.

---

## 2. Benchmark Output (Part 1 & 2)

- **Matrix Multiplication Result:** Success/Verified.
- **TLB Statistics:**
  - Total Lookups: 403

- Total Misses: 3
  - TLB Miss Rate: 0.007444
- 

### 3. Support for Different Page Sizes

- **Macro-based Design:** The implementation relies on constants defined in `my_vm.h` (`PGSIZE`, `MAX_MEMSIZE`, `PT_ENTRIES`, `PD_ENTRIES`).
  - **Bit Manipulation:** Address translation logic uses `OFFBITS`, `PDXSHIFT`, and `PTXSHIFT`. Changing the `PGSIZE` definition automatically adjusts the bit masks and shift values, allowing the code to support 8KB or larger pages without logic changes.
  - **Linear Scaling:** The bitmap sizes and `malloc` calculations dynamically adjust based on `MEMSIZE / PGSIZE` (defined as `NUM_PHYS_FRAMES`), ensuring the bitmap covers all frames regardless of page size.
- 

### 4. Possible Issues

- **Internal Fragmentation:** `n_malloc` allocates full pages even for small requests (e.g., 1 byte requests consume 4KB). This leads to wasted space within pages.
  - **Search Latency:** `get_next_avail` performs a linear search on the virtual bitmap. As memory fills up, finding a large contiguous block of free virtual memory becomes slower  $O(N)$ .
  - **TLB Thrashing:** Since the TLB is direct-mapped, alternating access between two pages that map to the same TLB index will cause constant evictions and a high miss rate.
- 

### 5. Extra Credit: 64-bit 4-Level Page Table

Refers to implementation in `my_vm64.c`.

#### Design Changes

- **4-Level Hierarchy:** Implements PML4 → PDPT → Page Directory → Page Table. This supports the larger virtual address space required for 64-bit systems.
- **Data Types:** All address variables, page table entries (`pte_t`), and directory entries (`pde_t`) are updated to `uint64_t`.
- **Virtual Address Space:** Supports 48-bit virtual addressing (`1ULL << 48`) as specified in the requirements.

#### VMA Management (Linked List)

- Unlike the 32-bit bitmap approach, the 64-bit implementation uses a linked list (`struct vm_area`) to manage virtual memory regions.
- **Nodes:** Each node tracks a `start` address and `size`.

- **Allocation:** `get_next_avail` traverses the sorted linked list to find a "gap" between allocations large enough to fit the requested number of pages. This is more efficient for the sparse 64-bit address space than a massive bitmap.

## Logic Adjustments

- **Translation:** The `translate` and `map_page` functions now cascade through 4 levels of translation. If any intermediate table (PML4, PDPT, PD) is missing during mapping, a new physical frame is allocated and zeroed out to create that table.
- **TLB:** The TLB struct was updated to store 64-bit VPNs and PFNs.
- **Thread Safety:** Maintains granular locking with `phys_bitmap_lock` (frames), `vm_lock` (VMA list), `pt_lock` (page tables), and `tlb_lock` (cache).