

Project 4 Report: RU File System (RUFS)

- Kelvin Ihezue (NetID: ki120)
- Bryan Shangguan (NetID: bys8)

Benchmark and Performance Analysis

Total Blocks Used

- Static Overhead (System Metadata):
 - Superblock: 1 Block (Block 0).
 - Inode Bitmap: 1 Block (Block 1).
 - Data Bitmap: 1 Block (Block 2).
 - Inode Region: Calculated as `(MAX_INUM * sizeof(struct inode)) / BLOCK_SIZE`.
 - With `MAX_INUM` = 1024 and `sizeof(struct inode)` ≈ 256 bytes, the inode region occupies roughly 64 Blocks.
 - Total Static Overhead: ~67 Blocks reserved immediately upon `rufs_mkfs`.
- Dynamic Usage (Benchmark Operations):
 - Root Directory: 1 Data Block (allocated during `rufs_mkfs`).
 - Test 1 (File Creation `/file`): Consumes 1 Inode (allocated from the reserved Inode Region).
 - Test 2 (File Write): The benchmark writes 16 iterations of 4096 bytes.
 - Total Data: 64KB.
 - Data Blocks Consumed: 16 Blocks.
 - Test 5 (Directory Creation `/files`): Consumes 1 Inode and 1 Data Block (for the new directory's entries).
 - Total Dynamic Data Blocks: 1 (Root) + 16 (File Data) + 1 (Dir Data) = 18 Blocks.
- Total Estimated Block Usage: 85 Blocks (67 Static + 18 Dynamic).

Implementation Strategy

Our implementation of RUFS follows a modular design centered around the FUSE API handlers in `rufs.c` and the block layer in `block.c`.

- Superblock & Initialization (`rufs_mkfs`, `rufs_init`):
 - We implemented `rufs_mkfs` to partition the "virtual disk" (flat file) into four distinct regions: Superblock, Bitmaps, Inode Table, and Data Blocks.
 - Critical validation occurs in `rufs_init`, where we check for the existence of `DISKFILE`. If missing, we format the disk; otherwise, we load the superblock into the global `sb` structure.
- Inode Operations (`readi`, `writei`):

- We implemented a mapping function that converts a logical inode number (`ino`) into a physical disk block and offset.
 - `readi` reads the specific block containing the requested inode into a buffer and extracts the specific inode struct.
 - `writei` performs a read-modify-write operation to update a single inode without corrupting neighboring inodes in the same block.
 - Bitmap Management:
 - We utilized the provided `set_bitmap`, `unset_bitmap`, and `get_bitmap` macros.
 - `get_avail_ino` and `get_avail_blkno` scan their respective blocks linearly to find the first zero bit, marking it as used immediately to prevent race conditions during allocation.
 - Directory Operations (`dir_find`, `dir_add`):
 - Lookup: `dir_find` iterates through the direct pointers of a directory inode. It reads each data block and casts it to an array of `struct dirent`, checking `valid` flags and comparing filenames.
 - Insertion: `dir_add` scans for an invalid (empty) `dirent` slot. If a block is full, it dynamically allocates a new block via `get_avail_blkno` and updates the directory inode's size and `direct_ptr`.
 - File Operations (`rufs_read`, `rufs_write`):
 - Read: Handles offsets by calculating the starting block index (`offset / BLOCK_SIZE`) and the intra-block offset (`offset % BLOCK_SIZE`). It copies data chunk-by-chunk into the user buffer.
 - Write: Similar to read, but handles dynamic allocation. If a write targets a block index that is currently `0` in the inode, we allocate a new block on the fly. We also handle file size updates in the inode metadata.
-

Difficulties and Issues Faced

Superblock Initialization Logic (Resolved)

- Issue: We initially encountered a critical bug where `rufs_mkfs` failed to create a valid root directory, resulting in "No such file or directory" errors during mounting.
- Cause: We were calling `writei(0, &root)` before initializing the global `sb` variable. The `writei` function relies on `sb.i_start_blk` to locate the inode region. Because `sb` was zeroed out (global default), `writei` wrote the root inode to Block 0, overwriting our Superblock and leaving the actual inode region empty.
- Solution: We reordered the code in `rufs_mkfs` to assign `sb = newsb` before calling `writei`. This ensured `writei` had the correct offsets.

Disk File Permissions (Resolved)

- Issue: We faced a `disk_open` failed: `Permission denied` error when restarting the program.
 - Cause: The `DISKFILE` was created in a previous session (possibly with different user privileges or flags), and `dev_open` attempted to open it with `O_RDWR` which failed.
 - Solution: We added a step to our workflow to explicitly `rm DISKFILE` before re-running the file system, ensuring `rufs_init` calls `rufs_mkfs` to create a fresh disk with the current user's ownership.
-

Unresolved Issues

- None.
 - We successfully resolved the critical logic bug in `rufs_mkfs` regarding the global superblock initialization.
 - We resolved the `writei` offset calculation errors.
 - The permission errors were resolved by cleaning the environment.
 - The benchmark cases (File Create, Small Write, Read, Directory Create) all pass successfully with the current codebase.
- Potential Future Improvements (If time permitted):
 - Indirect Pointers: Currently, the file system is limited to $16 * 4KB = 64KB$ per file because we only implemented direct pointers. Implementing indirect pointers would allow for much larger files.
 - Directory Entry Deletion: We did not implement `rufs_rmdir` or `rufs_unlink`. While not required for the assignment, this means the file system grows monotonically and cannot reclaim space.